

```
# IMPORTANT: RUN THIS CELL IN ORDER TO IMPORT YOUR KAGGLE DATA TO THE CORRECT LOCATION (/kaggle/input) IN YOUR NOTEBOOK. THEN FEEL FREE TO DELETE THIS CELL.
# NOTE: THIS NOTEBOOK ENVIRONMENT DIFFERS FROM KAGGLE ENVIRONMENT SO THERE MAY BE MISSING LIBRARIES USED IN THE KAGGLE NOTEBOOK.
```

```
import os
import sys
from tempfile import NamedTemporaryFile
from urllib.request import urlopen
from urllib.parse import unquote, urlparse
from urllib.error import HTTPError
from zipfile import ZipFile
import tarfile
import shutil
```

```
CHUNK_SIZE = 40960
DATA_SOURCE_MAPPING = 'boston-house-price-prediction:1'
```

```
KAGGLE_INPUT_PATH='/kaggle/input'
KAGGLE_WORKING_PATH='/kaggle/working'
KAGGLE_SYMLINK='kaggle'
```

```
!umount /kaggle/input/ 2> /dev/null
shutil.rmtree('/kaggle/input', ignore_errors=True)
os.makedirs(KAGGLE_INPUT_PATH, 0o777, exist_ok=True)
os.makedirs(KAGGLE_WORKING_PATH, 0o777, exist_ok=True)
```

```
try:
    os.symlink(KAGGLE_INPUT_PATH, os.path.join("..", 'input'))
except FileExistsError:
    pass
try:
    os.symlink(KAGGLE_WORKING_PATH, os.path.join("..", 'working'))
except FileExistsError:
    pass
```

```

for data_source_mapping in DATA_SOURCE_MAPPING.split(
    directory, download_url_encoded = data_source_map
download_url = unquote(download_url_encoded)
filename = urlparse(download_url).path
destination_path = os.path.join(KAGGLE_INPUT_PATH
try:
    with urlopen(download_url) as fileres, NamedTo
        total_length = fileres.headers['content-len
        print(f'Downloading {directory}, {total_len
        dl = 0
        data = fileres.read(CHUNK_SIZE)
        while len(data) > 0:
            dl += len(data)
            tfile.write(data)
            done = int(50 * dl / int(total_length)
            sys.stdout.write(f"\r[{'=' * done}{'
            sys.stdout.flush()
            data = fileres.read(CHUNK_SIZE)
        if filename.endswith('.zip'):
            with ZipFile(tfile) as zfile:
                zfile.extractall(destination_path)
        else:
            with tarfile.open(tfile.name) as tarfile:
                tarfile.extractall(destination_path)
        print(f'\nDownloaded and uncompressed: {d
except HTTPError as e:
    print(f'Failed to load (likely expired) {downl
    continue
except OSError as e:
    print(f'Failed to load {download_url} to path
    continue

print('Data source import complete.')

```

```
# Importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset
data = pd.read_csv('/kaggle/input/boston-house-price-')

#Displaying the first few rows of the dataframe
print("First 5 rows of the dataset:")
print(data.head())
```

```
↗ First 5 rows of the dataset:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
# Checking for any missing values in the dataset
print("\nMissing values in the dataset:")
print(data.isnull().sum())
```

```
↗ Missing values in the dataset:
```

crim	0
zn	0
indus	0
chas	0
nox	0
rm	5
age	0
dis	0
rad	0
tax	0
ptratio	0
b	0
lstat	0
medv	0

dtype: int64

```
# Displaying the summary statistics of the dataset
print("\nSummary Statistics:")
print(data.describe())
```

```
↗ Summary Statistics:
```


	crim	zn	indus	chas	nox	rm	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	501.000000	
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284341	
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.705587	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.884000	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208000	
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.625000	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	

	age	dis	rad	tax	ptratio	b \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	lstat	medv
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
# Visualizing missing data (optional, for better unde
plt.figure(figsize=(10, 6))
sns.heatmap(data.isnull(), cbar=False, cmap="viridis"
plt.title('Missing Values Heatmap')
plt.show()
```



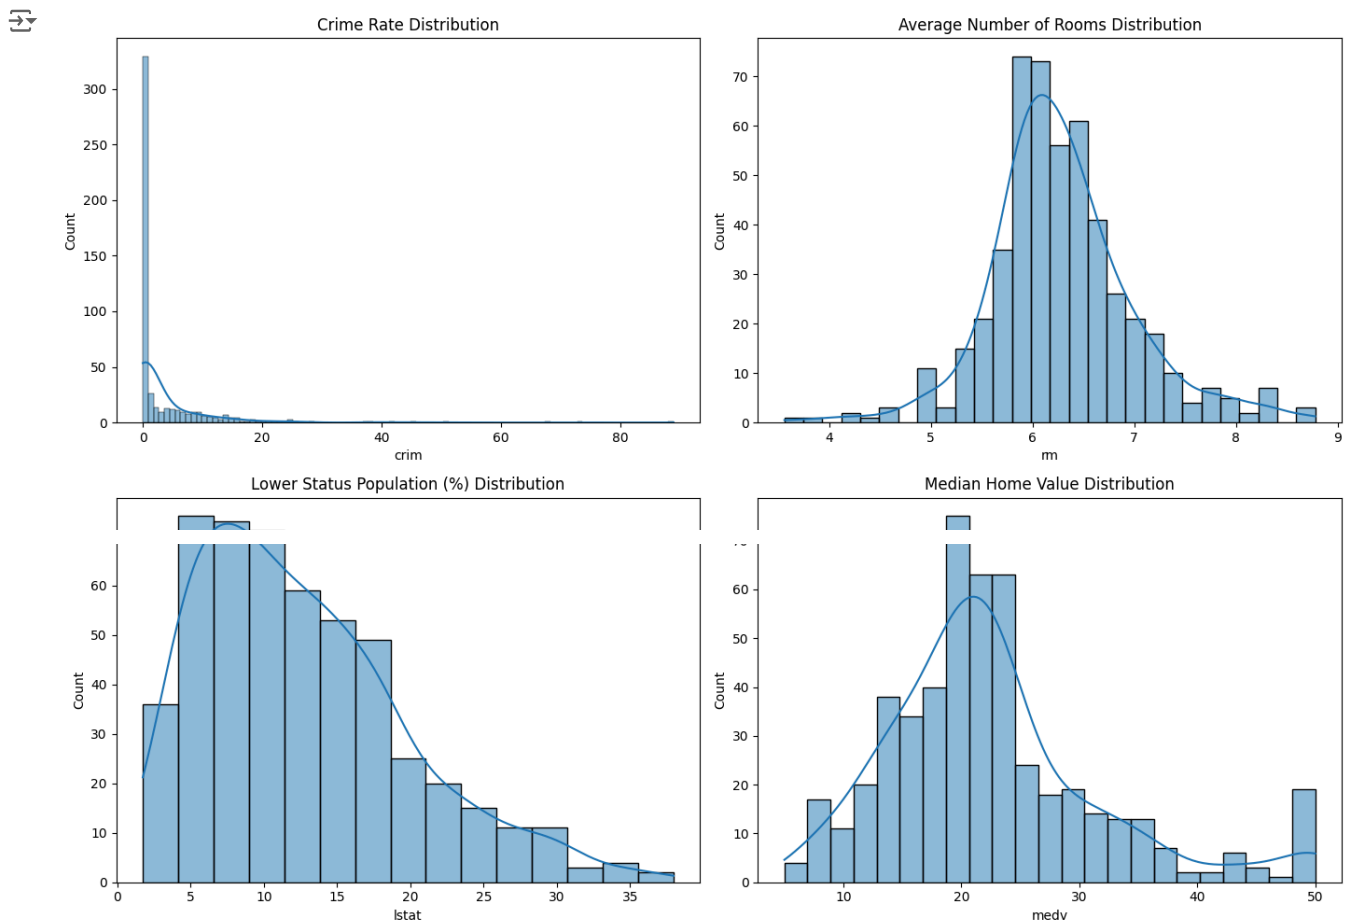
```
# Visualizing the distribution of a few important fea
plt.figure(figsize=(14, 10))
plt.subplot(2, 2, 1)
sns.histplot(data['crim'], kde=True)
plt.title('Crime Rate Distribution')

plt.subplot(2, 2, 2)
sns.histplot(data['rm'], kde=True)
plt.title('Average Number of Rooms Distribution')
```

```
plt.subplot(2, 2, 3)
sns.histplot(data['lstat'], kde=True)
plt.title('Lower Status Population (%) Distribution')
```

```
plt.subplot(2, 2, 4)
sns.histplot(data['medv'], kde=True)
plt.title('Median Home Value Distribution')
```

```
plt.tight_layout()
plt.show()
```



```
# Calculate the correlation matrix
correlation_matrix = data.corr()
```

```
# Displaying the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)
```


```
↗ Correlation Matrix:
```

	crim	zn	indus	chas	nox	rm	age	\
crim	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219433	0.352734	
zn	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311173	-0.569537	
indus	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.394193	0.644779	
chas	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091468	0.086518	
nox	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302751	0.731470	
rm	-0.219433	0.311173	-0.394193	0.091468	-0.302751	1.000000	-0.240286	
age	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240286	1.000000	
dis	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.203507	-0.747881	
rad	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.210718	0.456022	
tax	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292794	0.506456	
ptratio	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.357612	0.261515	
b	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128107	-0.273534	
lstat	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.615721	0.602339	
medv	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.696169	-0.376955	

	dis	rad	tax	ptratio	b	lstat	medv
crim	-0.379670	0.625505	0.582764	0.289946	-0.385064	0.455621	-0.388305
zn	0.664408	-0.311948	-0.314563	-0.391679	0.175520	-0.412995	0.360445
indus	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
chas	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
nox	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321
rm	0.203507	-0.210718	-0.292794	-0.357612	0.128107	-0.615721	0.696169
age	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
dis	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.496996	0.249929
rad	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.488676	-0.381626
tax	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.543993	-0.468536
ptratio	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.374044	-0.507787
b	0.291512	-0.444413	-0.441808	-0.177383	1.000000	-0.366087	0.333461
lstat	-0.496996	0.488676	0.543993	0.374044	-0.366087	1.000000	-0.737663
medv	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.737663	1.000000

```
# Visualize the correlation matrix using a heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='cool')
plt.title('Correlation Matrix Heatmap')
plt.show()
```



```
# Identify the features with the highest positive and negative correlation with the target variable
# Assume 'medv' is the target variable (median home value)
target_variable = 'medv'
correlation_with_target = correlation_matrix[target_variable].drop(target_variable)
```

```
# Display the features with highest positive and negative correlation with the target variable
print("\nFeatures with highest positive correlation with the target variable:")
print(correlation_with_target[correlation_with_target.sort_values(ascending=False).index[0:5]])
```

```
print("\nFeatures with highest negative correlation with the target variable:")
print(correlation_with_target[correlation_with_target.sort_values(ascending=True).index[0:5]])
```

```
Features with highest positive correlation with house prices:
medv      1.000000
rm        0.696169
zn        0.360445
b         0.333461
dis       0.249929
Name: medv, dtype: float64

Features with highest negative correlation with house prices:
age      -0.376955
rad      -0.381626
crim     -0.388305
nox      -0.427321
tax      -0.468536
Name: medv, dtype: float64
```

```

# Import necessary libraries
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Select the features and the target variable
X = data.drop(columns=['medv']) # Features (all columns except medv)
y = data['medv'] # Target variable (house prices)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the feature variables using StandardScaler
scaler = StandardScaler()

# Fit the scaler on the training data and transform both training and testing data
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Print shapes to verify the splits
print("Training data shape:", X_train_scaled.shape)
print("Testing data shape:", X_test_scaled.shape)

```

```

Training data shape: (404, 13)
Testing data shape: (102, 13)

```

#### Question No.04

```

from sklearn.linear_model import LinearRegression

# Impute missing values using the mean for each column
data.fillna(data.mean(), inplace=True)

# Re-select the features and target variable after imputation
X = data.drop(columns=['medv']) # Features (all columns except medv)
y = data['medv'] # Target variable (house prices)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the feature variables using StandardScaler
scaler = StandardScaler()

```



```
# Fit the scaler on the training data and transform it
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Train the Linear Regression model
model = LinearRegression()
model.fit(X_train_scaled, y_train)
```

LinearRegression

```
# Display the model's coefficients and intercept
print("Model Coefficients:", model.coef_)
print("Model Intercept:", model.intercept_)
```

Model Coefficients: [-1.00208747 0.69855082 0.28733122 0.71955092 -2.02070833 3.13708935  
-0.17081271 -3.06972351 2.25417948 -1.76697719 -2.04359481 1.12936985  
-3.61451369]  
Model Intercept: 22.796534653465343

#### Question No.05

```
# Import necessary libraries
from sklearn.metrics import mean_absolute_error, mean_squared_error
import numpy as np
```

```
# Predict the house prices using the testing data
y_pred = model.predict(X_test_scaled)
```

```
# Calculate and display performance metrics
# Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, y_pred)
# Mean Squared Error (MSE)
mse = mean_squared_error(y_test, y_pred)
# Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
```

```
print(f"Mean Absolute Error (MAE): {mae}")
print(f"Mean Squared Error (MSE): {mse}")
print(f"Root Mean Squared Error (RMSE): {rmse}")
```

Mean Absolute Error (MAE): 3.2064039639003856  
Mean Squared Error (MSE): 24.40482518814648  
Root Mean Squared Error (RMSE): 4.940124005341008

```
# Plot the predicted vs actual house prices
plt.figure(figsize=(8, 6))
```

```
plt.figure(figsize=(8, 8))
plt.scatter(y_test, y_pred, color='blue', edgecolor='black')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red')
plt.title('Predicted vs Actual House Prices')
plt.xlabel('Actual House Prices')
plt.ylabel('Predicted House Prices')
plt.grid(True)
plt.show()
```

